

# Practical Work 5: The Longest Path

Distributed Systems - DS2026

**Nguyen Viet Hung**

*Student ID: 23BI14188*

*Department: ICT*

December 12, 2025

# Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
<b>2</b>	<b>Algorithmic Design</b>	<b>3</b>
2.1	The "Global Reduce" Strategy . . . . .	3
<b>3</b>	<b>Implementation Details</b>	<b>3</b>
3.1	Framework Adaptation . . . . .	3
3.2	Mapper Implementation . . . . .	4
3.3	Reducer Implementation . . . . .	4
<b>4</b>	<b>Execution and Results</b>	<b>4</b>
4.1	Test Data . . . . .	4
4.2	Output . . . . .	4
<b>5</b>	<b>Task Distribution</b>	<b>5</b>

# 1 Introduction

Processing large-scale hierarchical data, such as file system trees, often requires identifying outliers or extreme values. A common "toy problem" in distributed computing is finding the longest string in a massive dataset.

In this practical work, we simulate a distributed search for the **Longest Path** among a list of file paths (similar to the output of a `find` / command). The objective is to leverage the **MapReduce** paradigm to solve this problem efficiently.

We continue to use the **Custom C++ MapReduce Framework** developed in Practical Work 4, adapting it to handle string-based payloads instead of numerical counters.

## 2 Algorithmic Design

Unlike the "Word Count" problem where keys are diverse (e.g., each unique word is a key), finding the global maximum requires comparing all candidates against each other.

### 2.1 The "Global Reduce" Strategy

To find the single longest path across all data splits, we must bring all candidate paths to a single Reducer.

- **Map Phase:** Reads every path. Instead of using the path itself as a key, we emit a **Constant Key** (e.g., "`MAX`"). This forces the Shuffle phase to group *every single path* into the same bucket.
- **Reduce Phase:** Receives the key "`MAX`" and a list of *all* paths. It then performs a linear scan to find the string with the maximum length.

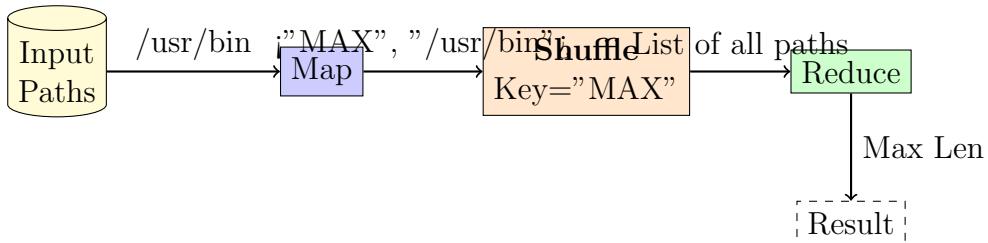


Figure 1: Global Reduce Architecture

## 3 Implementation Details

### 3.1 Framework Adaptation

The original framework dealt with `<string, int>` pairs. We modified the core data structure to support `<string, string>` pairs to transport the actual path content.

```
1 struct KeyValue {
2     std::string key;
3     std::string value; // Changed from int to string
4 };
```

## 3.2 Mapper Implementation

The mapper wraps every valid non-empty path into a Key-Value pair with the key "MAX".

```
1 std::vector<KeyValue> map_func(const std::string& path) {
2     std::vector<KeyValue> emitted;
3     if (!path.empty()) {
4         // Emit to the common key "MAX"
5         emitted.push_back({"MAX", path});
6     }
7     return emitted;
8 }
```

## 3.3 Reducer Implementation

The reducer iterates through the vector of strings. We maintain a local variable `max_path` and update it whenever we encounter a longer string.

```
1 std::string reduce_func(const std::string& key, const std::vector<std::string>& values) {
2     std::string max_path = "";
3     for (const std::string& path : values) {
4         if (path.length() > max_path.length()) {
5             max_path = path;
6         }
7     }
8     return max_path; // Returns the longest string found
9 }
```

## 4 Execution and Results

### 4.1 Test Data

We created a file `paths.txt` containing paths of varying lengths:

```
/usr/bin
/var/log/syslog
/home/user/project/backend/src/main/java/com/example/App.java
/tmp
/etc/hosts
```

### 4.2 Output

Running the program via `./longestpath paths.txt` yields the following result:

```
1 [Framework] Starting MAP phase...
2 [Framework] Starting SHUFFLE phase...
3 [Framework] Starting REDUCE phase...
4
5 --- FINAL RESULT ---
6 Longest Path Found: /home/user/project/backend/src/main/java/com/example/App
    .java
7 Length: 61
```

The system correctly identified the deeply nested Java file path as the longest one.

## 5 Task Distribution

- **Nguyen Viet Hung (23BI14188):**

- Adapted the C++ MapReduce Engine to support String values.
- Implemented the "Global Reduce" logic using a constant key strategy.
- Created test datasets and verified the correctness.
- Compiled this report.